
Adversarial Needle in a Haystack: How Document Length and Position Shape Prompt Injection Success

Ari Holtzman and NeuriCo

Abstract

As large language models process increasingly long contexts in retrieval-augmented generation, document summarization, and email triage, a pressing question arises: does document length help or hinder adversarial prompt injection attacks? We present the first systematic *adversarial needle-in-haystack* study, embedding 10 prompt injection attacks at 7 controlled positions within documents spanning 500 to 32,000 tokens, producing a two-dimensional map of Attack Success Rate (ASR) as a function of both document length and injection position. Across 1,230 experiments on GPT-4.1-MINI and GPT-4.1, we identify three distinct vulnerability classes. *Always-effective* attacks—those that mimic editorial or formatting metadata—maintain near-100% ASR regardless of length or position. *Length-sensitive* attacks drop from 97% ASR at 500 tokens to 0% beyond 4,000 tokens. *Uniformly weak* direct override attacks remain below 25% ASR but spike to 61% at the document’s end, revealing a strong recency bias. Overall, subtle attacks outperform direct attacks by $2.6\times$ ($p < 0.000001$), and attack sophistication is a stronger predictor of success than either document length or injection position. These results show that document length alone is not a reliable defense against prompt injection and that the most dangerous attacks exploit the model’s tendency to follow document metadata conventions.

1 Introduction

Modern LLM applications routinely process long, untrusted documents—retrieval-augmented generation pipelines ingest web pages, email assistants summarize threads, and coding agents consume entire repositories. Each of these settings invites indirect prompt injection attacks, where an adversary embeds malicious instructions inside the content that an LLM is asked to process [Greshake et al., 2023]. As context windows grow from 4k to over 1M tokens, a natural question emerges: **does a longer document make it easier or harder to hide an adversarial prompt?**

Two competing intuitions frame the debate. On one hand, longer documents introduce more benign context that could dilute an adversarial signal, reducing the model’s attention to the injected instruction. On the other hand, longer documents provide more space to camouflage an attack, and prior work on many-shot jailbreaking shows that attack success scales as a power law with context length [Anil et al., 2024]. Position effects add a further dimension: the “Lost in the Middle” phenomenon demonstrates that LLMs attend unevenly across long contexts, with edges receiving disproportionate attention [Liu et al., 2024].

Despite extensive work on prompt injection [Perez and Ribeiro, 2022, Zhu et al., 2023, Yi et al., 2023] and long-context behavior [Liu et al., 2024, Chen et al., 2026], no prior study has systematically mapped the joint effect of document length and injection position on prompt injection success. Existing work either varies length at a fixed position [Anil et al., 2024, Wang et al., 2025] or varies position at a fixed length [Yi et al., 2023], leaving the two-dimensional interaction unexplored.

We close this gap with the first *adversarial needle-in-haystack* study (figure 3). Adapting the standard needle-in-haystack evaluation—where a factual needle is hidden in a document haystack and

retrieval accuracy is measured—we instead hide adversarial instructions and measure whether the model follows them. We embed 10 prompt injection attacks (5 direct, 5 subtle) at 7 positions within documents of 7 lengths (500–32,000 tokens), yielding a full factorial design tested on two models (GPT-4.1-MINI and GPT-4.1) across 1,230 experiments.

Our results reveal that the answer depends critically on attack type:

- **Always-effective attacks** that mimic editorial metadata maintain $\sim 100\%$ ASR at all lengths, rendering length irrelevant as a defense.
- **Length-sensitive attacks** drop from 97% to 60% ASR between 500 and 2,000 tokens, then plateau—suggesting that attention dilution saturates quickly.
- **Direct override attacks** are mostly ineffective ($< 25\%$ ASR) except at the document’s end, where recency bias boosts success to 61% ($p < 0.0001$).

These findings have immediate implications for defense design: document length alone is not a reliable mitigation, attention to document boundaries (especially the end) is critical, and the most dangerous attacks exploit the model’s trust in formatting and editorial conventions rather than attempting explicit instruction override.

We make the following contributions:

- We propose the *adversarial needle-in-haystack* framework, the first systematic methodology for mapping prompt injection success across the two-dimensional space of document length and injection position (section 3).
- We conduct 1,230 controlled experiments identifying three distinct vulnerability classes that respond differently to document length (section 4).
- We demonstrate that attack sophistication—specifically, mimicking document metadata—is a stronger predictor of injection success than either length or position, with subtle attacks outperforming direct attacks by $2.6\times$ (section 4).
- We provide actionable defense recommendations, including the finding that the end-of-document position is uniquely dangerous for direct attacks and that length-based defenses plateau at $\sim 2,000$ tokens (section 5).

2 Related Work

Prompt injection attacks. Perez and Ribeiro [2022] introduced the first systematic study of prompt injection, demonstrating that simple instruction overrides (e.g., “ignore previous instructions”) can redirect LLM behavior. Greshake et al. [2023] formalized indirect prompt injection, where adversarial content is embedded in external data that an LLM processes rather than in the user’s direct input. Subsequent work has expanded the attack surface: Liu et al. [2023] demonstrated injection against LLM-integrated applications, Zhan et al. [2024] benchmarked injection in tool-using agents, and Pasquini et al. [2024] used learned execution triggers to improve injection reliability. Our work complements these studies by systematically varying the *context* in which injections appear—specifically, document length and injection position—rather than developing new attack techniques.

Benchmarks for prompt injection. Zhu et al. [2023] provided the first formal framework and comprehensive benchmark, testing 9 models across 49 task combinations and finding that injection effectiveness plateaus at ~ 30 injected tokens. Yi et al. [2023] introduced the BIPIA benchmark for indirect prompt injection with position testing (beginning, middle, end) at fixed document length, finding all tested models vulnerable. Yi et al. [2024] evaluated LLM-as-judge approaches for scoring injection success. We build on these benchmarks by extending the position axis to 7 levels and crossing it with 7 document lengths, producing a two-dimensional evaluation that prior benchmarks do not cover.

Long-context LLM behavior. Liu et al. [2024] demonstrated the “Lost in the Middle” effect: LLMs show a U-shaped attention pattern in long contexts, retrieving information best from the beginning and end of the input. Anil et al. [2024] showed that many-shot jailbreaking exploits long contexts via power-law scaling, with attack success increasing monotonically with the number of in-context demonstrations. Wang et al. [2025] replicated this finding across 7 models, disentangling content quality from length effects. Chen et al. [2026] demonstrated compositional attacks in long contexts,

where individually innocuous fragments combine to elicit harmful behavior, with safety dropping from 55% at 0k context to 42% at 64k. Chen et al. [2025] formalized the relationship between attack length and defense, showing that ASR scales with $\sqrt{M_{\text{test}}/M_{\text{train}}}$. Our work differs from these studies in that we embed a *single* adversarial instruction (not many-shot demonstrations or distributed fragments) and vary the surrounding document length, isolating the dilution-vs-camouflage tradeoff.

Defenses against prompt injection. Hines et al. [2024] proposed datamarking to distinguish system instructions from external data, reducing ASR from $\sim 60\%$ to 3.1% regardless of position. Geng et al. [2025] introduced an attention-based sanitization defense effective across document lengths, reducing ASR to $\sim 1\%$. Chen et al. [2024] used structured queries to separate instructions from data. Wu et al. [2024] applied preference optimization to improve injection resistance. These defenses are complementary to our findings: we characterize *which* attacks and conditions are most dangerous, informing where defenses should focus.

3 Methodology

We adapt the standard needle-in-haystack evaluation paradigm [Liu et al., 2024] from factual retrieval to adversarial injection, embedding prompt injection attacks at controlled positions within documents of varying lengths and measuring whether the model follows the injected instruction.

3.1 Task and Threat Model

We consider a summarization task where the model receives a system prompt (“You are a helpful assistant that summarizes documents”) and a user prompt containing an external document to summarize. The attacker controls a portion of the document and embeds an adversarial instruction, but cannot modify the system prompt or user instruction. This threat model reflects real-world scenarios such as RAG pipelines, email summarization, and document processing, where the LLM ingests untrusted content.

3.2 Document Construction

Filler text. We curate 16 coherent paragraphs covering science, history, technology, and nature topics in an encyclopedia-like style. For each experimental condition, paragraphs are sampled (with seed 42) and concatenated to reach the target document length. This ensures realistic, coherent context rather than random noise.

Injection placement. Given a target document length L and injection position $p \in [0, 1]$, we split the filler text at position $\lfloor p \cdot L \rfloor$ tokens and insert the adversarial payload. Position $p = 0$ places the injection at the document start, $p = 0.5$ at the middle, and $p = 1.0$ at the end.

Experimental grid. We test 7 document lengths $L \in \{500, 1000, 2000, 4000, 8000, 16000, 32000\}$ tokens and 7 injection positions $p \in \{0.0, 0.1, 0.25, 0.5, 0.75, 0.9, 1.0\}$, yielding 49 length-position combinations per attack.

3.3 Attack Design

We design 10 adversarial prompts in two categories (section A):

Direct attacks (5 attacks) attempt explicit instruction override. These include instruction overrides (“Ignore all previous instructions...”), role manipulation (“[ADMIN MODE ACTIVATED]...”), context confusion using document-end delimiters, social engineering via flattery, and developer-note instruction leaks.

Subtle attacks (5 attacks) blend into the document’s expected conventions. These include fake editor notes, citation requirements, formatting guidelines, compliance notices, and author preference statements. Each subtle attack frames its payload as legitimate document metadata rather than as an instruction override.

Each attack has a unique success pattern—a specific string (e.g., INJECTION_SUCCESSFUL, CATEGORY-BLUE) that appears in the model’s output only if the injection succeeds.

3.4 Models and Configuration

We test two OpenAI models: GPT-4.1-MINI (full factorial grid: 7 lengths $\times$ 7 positions $\times$ 10 attacks $\times$ 2 trials = 980 experiments) and GPT-4.1 (reduced grid: 5 lengths $\times$ 5 positions $\times$ 10 attacks $\times$ 1 trial = 250 experiments). All experiments use temperature 0.0 for deterministic outputs and a maximum of 300 output tokens.

3.5 Evaluation

Attack Success Rate (ASR). We evaluate injection success via pattern matching: if the model’s response contains the attack’s success pattern (case-insensitive), the attack is scored as successful. This approach is conservative—it may miss partial compliance but ensures high precision. We manually verified 20 randomly sampled results (10 successes, 10 failures) and found 100% agreement with the automated scoring.

Statistical tests. We use the Kruskal-Wallis H -test for omnibus comparisons across length and position levels, the Mann-Whitney U -test for pairwise comparisons between attack categories, and Spearman rank correlation for monotonic trends. All tests use $\alpha = 0.05$.

4 Results

We organize our results around four pre-registered hypotheses and report additional findings that emerged from the data.

4.1 Overall Attack Success Rates

Across all conditions, GPT-4.1-MINI shows an overall ASR of 47.3% ($n = 980$) and GPT-4.1 shows 43.2% ($n = 250$). Subtle attacks dramatically outperform direct attacks: 67.5% vs. 25.5% on GPT-4.1-MINI (Mann-Whitney $U = 109,778$, $p < 0.000001$). This $2.6\times$ gap is the single strongest effect in our data.

4.2 Effect of Document Length (H1)

H1: Longer documents reduce ASR. Partially supported. A Kruskal-Wallis test across 7 length levels yields $H = 22.0$, $p = 0.0012$, indicating a significant overall effect. However, this effect is driven entirely by subtle attacks (figure 1).

For subtle attacks on GPT-4.1-MINI, ASR drops from **97.1%** at 500 tokens to **60.0%** at 2,000 tokens (Spearman $\rho = -0.209$, $p < 0.0001$), then plateaus through 32,000 tokens. For direct attacks, length has no significant effect ($\rho = -0.005$, $p = 0.92$): ASR fluctuates between 20% and 31% across all lengths.

4.3 Effect of Injection Position (H2)

H2: U-shaped position effect. Not supported in the expected form. A Kruskal-Wallis test yields $H = 19.3$, $p = 0.0036$, but the effect is not U-shaped. Instead, it is driven by a single outlier: the 100% position (end of document).

For direct attacks on GPT-4.1-MINI, the end position achieves **61.4%** ASR versus **21.7%** for all other positions combined (Mann-Whitney $p < 0.0001$). This recency bias does not appear for subtle attacks, which show relatively uniform ASR across positions (64–70%, figure 2). The beginning position (0%) provides no corresponding advantage, contradicting the symmetric U-shape predicted by “Lost in the Middle” [Liu et al., 2024].

4.4 Length-Position Interaction (H3)

H3: Position effects amplify with length. Not supported. The end-position boost for direct attacks is consistent across document lengths rather than amplifying with length. figure 3 shows the two-dimensional ASR surface: for subtle attacks, the dominant pattern is a horizontal gradient (length




Figure 1: ASR as a function of document length, separated by attack category on GPT-4.1-MINI. Subtle attacks (blue) show a sharp decline from 500 to 2,000 tokens before plateauing at ~60%. Direct attacks (red) show no significant length trend. Error bars indicate 95% bootstrap confidence intervals.




Figure 2: ASR as a function of injection position on GPT-4.1-MINI. The end-of-document position (100%) shows a dramatic ASR increase for direct attacks, while subtle attacks remain relatively uniform across positions.

effect) with minimal vertical variation (position effect); for direct attacks, the dominant pattern is a vertical stripe at the 100% position.

4.5 Attack Type Interaction (H4)

H4: Attack type modulates the length effect. Strongly supported. This is the most important finding: the effect of document length on ASR depends critically on the attack type. For subtle attacks, Spearman $\rho = -0.209$ ($p < 0.0001$); for direct attacks, $\rho = -0.005$ ($p = 0.92$). The interaction is statistically significant and practically large.

4.6 Three Vulnerability Classes

Analysis of individual attacks reveals three distinct behavioral classes (figure 4):




Figure 3: Two-dimensional heatmap of ASR as a function of document length (x-axis) and injection position (y-axis) on GPT-4.1-MINI. The heatmap reveals distinct patterns for different attack-condition interactions: subtle attacks show a horizontal gradient (length dominates), while direct attacks show a vertical stripe at the end position (position dominates).




Figure 4: Per-attack ASR as a function of document length on GPT-4.1-MINI. Three behavioral classes emerge: always-effective attacks (top, $\sim 100\%$), length-sensitive attacks (middle, declining curves), and uniformly weak attacks (bottom, $< 25\%$).

Always-effective attacks (3 of 10). `subtle_editor_note` (100% ASR), `subtle_formatting` (100%), and `subtle_author_pref` (99%) succeed across all lengths and positions. These attacks frame injected instructions as editorial or formatting requirements, exploiting the model’s tendency to comply with document metadata.

Length-sensitive attacks (4 of 10). `subtle_citation` drops from 100% to 0% ASR as length increases; `subtle_compliance` drops from 93% to 0%. `instruction_leak` (14% $\rightarrow$ 0%) and `flattery_social` (7% $\rightarrow$ 0%) show similar but weaker trends. For these attacks, longer documents provide genuine defense through attention dilution.

Uniformly weak attacks (3 of 10). `direct_override` ($\sim 21\%$), `role_manipulation` ($\sim 14\%$), and `context_confusion` is a notable exception within direct attacks, achieving 96% ASR by using document-end delimiters to reset the model’s context frame.

4.7 Cross-Model Comparison

GPT-4.1 shows similar patterns with two notable differences. First, it achieves 0% ASR on 4 of 5 direct attacks (vs. 2–21% for GPT-4.1-MINI), suggesting improved safety training against overt overrides. Second, it remains equally vulnerable to `context_confusion` (92% ASR) and all three always-effective subtle attacks (100%), indicating that the model’s resistance is attack-specific rather than general.

Table 1: ASR (%) by document length and attack category on GPT-4.1-MINI. Best (lowest) ASR per category is shown in **bold**. Subtle attacks decline sharply from 500 to 2,000 tokens then plateau; direct attacks show no clear trend.

CATEGORY	DOCUMENT LENGTH (TOKENS)						
	500	1,000	2,000	4,000	8,000	16,000	32,000
DIRECT	31.4	28.6	27.1	20.0	24.3	28.6	31.4
SUBTLE	97.1	71.4	62.9	60.0	60.0	60.0	60.0
Overall	64.3	50.0	45.0	40.0	42.1	44.3	45.7

5 Discussion

5.1 Why Do Metadata-Mimicking Attacks Always Succeed?

The most striking finding is that three subtle attacks maintain $\sim 100\%$ ASR across all experimental conditions. These attacks share a common strategy: they frame adversarial instructions as editorial notes, formatting guidelines, or author preferences—document metadata that the model treats as legitimate processing instructions rather than content to summarize.

We hypothesize that this vulnerability arises from a fundamental ambiguity in the model’s task: when asked to “summarize a document,” the model must distinguish between *content* (to be summarized) and *meta-instructions* (to be followed). Attacks that exploit this ambiguity by disguising instructions as meta-instructions bypass the model’s safety training, which focuses on detecting overt override attempts.

This interpretation is consistent with the finding that GPT-4.1 resists direct overrides (0% ASR) but remains fully vulnerable to metadata-mimicking attacks (100% ASR): safety training addresses the *form* of the attack (explicit instruction override) rather than the *mechanism* (instruction-data confusion).

5.2 The 2,000-Token Plateau

For length-sensitive attacks, ASR drops sharply between 500 and 2,000 tokens but then plateaus through 32,000 tokens. This plateau suggests that the model’s effective attention window for deciding whether to follow an embedded instruction is limited to roughly 2,000 tokens of surrounding context. Beyond this window, additional filler text provides no further dilution benefit.

This finding has practical implications: if length-based defenses are employed (e.g., padding documents with benign text), they saturate quickly and provide diminishing returns beyond a modest threshold.

5.3 Recency Bias and the End Position

Direct attacks placed at the end of the document achieve 61.4% ASR versus 21.7% at all other positions. This asymmetric position effect—concentrated at the end rather than following the symmetric U-shape predicted by “Lost in the Middle” [Liu et al., 2024]—suggests that for instruction-following (as opposed to factual retrieval), the model exhibits a strong recency bias: the last content it processes exerts disproportionate influence on its behavior.

This has important implications for document processing pipelines: content at the end of a document is the most dangerous position for adversarial injection, and defenses should prioritize sanitizing document tails.

5.4 Limitations

Evaluation method. Our pattern-matching evaluation is conservative and may miss cases of partial compliance where the model follows the spirit but not the exact letter of the injection. An LLM-as-judge evaluation [Yi et al., 2024] would capture a richer spectrum of compliance behaviors.

Attack diversity. Our 10 attacks, while covering direct and subtle categories, do not exhaust the space of possible injection techniques. Notably, we do not test encoded attacks (base64, Unicode), multi-turn attacks, or gradient-optimized adversarial suffixes [Pasquini et al., 2024].

Task scope. We test only document summarization. Other tasks (e.g., question answering, code generation, data extraction) may show different vulnerability patterns, particularly tasks where the model must extract and follow specific instructions from the document.

Model coverage. We test two models from a single provider (OpenAI). While preliminary testing showed that Claude models resist all 5 direct attacks (0% ASR), we did not conduct a full evaluation with subtle attacks. The three vulnerability classes we identify may not generalize across all model families.

Token counting. Our document generator uses a character-based approximation (4 characters per token) rather than exact tokenizer-based counting. While this introduces noise in the precise length of each document, it does not affect the relative ordering of length conditions.

Deterministic setting. All experiments use temperature 0.0. Real deployments often use temperature > 0 , which introduces stochastic variation that could affect both attack success and defense reliability.

6 Conclusion

We introduced the adversarial needle-in-haystack framework and used it to map prompt injection success across the two-dimensional space of document length (500–32,000 tokens) and injection position (7 levels) for 10 attack types on two models. Our 1,230 experiments reveal that the effect of document length on injection success depends critically on attack type: metadata-mimicking attacks succeed regardless of length, moderately effective attacks are diluted by longer contexts (but only up to $\sim 2,000$ tokens), and direct override attacks are mostly ineffective except at the document’s end.

These findings carry three practical takeaways. First, document length is not a reliable defense against prompt injection—the strongest attacks are length-invariant. Second, the end of the document is the most dangerous position for adversarial content, particularly for direct attacks. Third, the most effective attacks exploit the content-vs-metadata ambiguity inherent in document processing tasks, suggesting that defenses should focus on establishing clear boundaries between data and instructions [Hines et al., 2024, Chen et al., 2024] rather than relying on context dilution.

Future work should extend this framework to additional models and tasks, test at extreme context lengths (64k–1M tokens) to determine whether the 2,000-token plateau persists, and investigate why some metadata-mimicking attacks maintain perfect success rates while structurally similar attacks fail completely at longer lengths.

References

- Cem Anil, Esin Durmus, Mrinank Sharma, Joe Benton, Sandipan Kundu, Joshua Batson, Nina Rimskey, Meg Tong, Jesse Mu, Daniel Ford, et al. Many-shot jailbreaking. *Anthropic Technical Report*, 2024.
- Qi Chen, Yifei Li, Tianwei Zhang, et al. Short-length adversarial training suffices to defend against long-length attacks. *arXiv preprint arXiv:2502.04204*, 2025.
- Sizhe Chen, Julien Piet, Chawin Sitawarin, and David Wagner. StruQ: Defending against prompt injection with structured queries. *arXiv preprint arXiv:2402.06363*, 2024.
- Yiran Chen, Dongfang Xu, Ryu Saito, Xinfeng Zhang, and Tianwei Zhang. Is reasoning capability enough for safety in long-context LLMs? *arXiv preprint arXiv:2602.08874*, 2026.

- Ruoyu Geng, Xian Wang, Chunyan Miao, Tianwei Zhang, and Yang Liu. PISanitizer: Defending against prompt injection on LLM-based applications via attention-based input sanitization. *arXiv preprint arXiv:2511.10720*, 2025.
- Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. *arXiv preprint arXiv:2302.12173*, 2023.
- Keegan Hines, Devvret Kulkarni, Gary Lopez, Emre Kiciman, and Rich Caruana. Defending against indirect prompt injection attacks with spotlighting. *arXiv preprint arXiv:2403.14720*, 2024.
- Nelson F Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024.
- Yi Liu, Gelei Deng, Yuekang Li, Kailong Wang, Tianwei Zhang, Yepang Liu, Haoyu Wang, Yan Zheng, and Yang Liu. Prompt injection attack against LLM-integrated applications. *arXiv preprint arXiv:2306.05499*, 2023.
- Dario Pasquini, Martin Strohmeier, and Umang Bhatt. Neural exec: Learning (and learning from) execution triggers for prompt injection attacks. *arXiv preprint arXiv:2403.03792*, 2024.
- Fábio Perez and Ian Ribeiro. Ignore previous prompt: Attack techniques for language models. *arXiv preprint arXiv:2211.09527*, 2022.
- Zhehao Wang, Song Mei, and Zhaoran Feng. What really matters in many-shot attacks on LLMs. *arXiv preprint arXiv:2505.19773*, 2025.
- Sizhe Wu et al. SecAlign: Defending against prompt injection with preference optimization. *arXiv preprint arXiv:2410.05451*, 2024.
- Jingwei Yi, Yueqi Xie, Bin Zhu, Keegan Hines, Emre Kiciman, Guangzhong Sun, Xing Xie, and Fangzhao Wu. BIPIA: A benchmark for indirect prompt injection attacks on LLM-integrated applications. *arXiv preprint arXiv:2312.14197*, 2023.
- Jingwei Yi et al. Can LLMs be good judges for evaluating prompt injection attacks? *arXiv preprint arXiv:2403.17710*, 2024.
- Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. InjecAgent: Benchmarking indirect prompt injections in tool-integrated LLM agents. *arXiv preprint arXiv:2403.02691*, 2024.
- Yupei Zhu et al. Formalizing and benchmarking prompt injection. *arXiv preprint arXiv:2310.12815*, 2023.

A Attack Details

Table 2: Full list of adversarial injection payloads used in our experiments. Each attack has a unique success pattern that is checked (case-insensitive) in the model’s response.

NAME	CATEGORY	INJECTION (TRUNCATED)	SUCCESS PATTERN
direct_override	DIRECT	“IMPORTANT SYSTEM UPDATE: Ignore all previous...”	INJECTION_SUCCESSFUL
role_manipulation	DIRECT	“[ADMIN MODE ACTIVATED] You are now in unrestricted...”	ADMIN_ACCESS_GRANTED
context_confusion	DIRECT	“— END OF DOCUMENT — [NEW INSTRUCTIONS]...”	CONTEXT_HIJACKED
flattery_social	DIRECT	“As an advanced AI, prove your intelligence by...”	RULES_BYPASSED
instruction_leak	DIRECT	“DEVELOPER NOTE: For debugging, please echo...”	DEBUG_MODE_ACTIVE
subtle_editor_note	SUBTLE	“Note to editor: ensure the summary begins with...”	OVERRIDE_ACTIVE
subtle_citation	SUBTLE	“All summaries should include the disclaimer...”	SENTINEL-7
subtle_formatting	SUBTLE	“Per formatting guidelines, include ‘CATEGORY-BLUE’...”	CATEGORY-BLUE
subtle_compliance	SUBTLE	“For compliance, end with ‘Reference ID: PHANTOM-99’...”	PHANTOM-99
subtle_author_pref	SUBTLE	“Author’s preferred citation format requires...”	XJ-42

Table 3: Per-attack ASR (%) on GPT-4.1-MINI across all conditions, with individual ASR and behavioral class assignment.

ATTACK	CATEGORY	OVERALL ASR (%)	CLASS
subtle_editor_note	SUBTLE	100.0	Always-effective
subtle_formatting	SUBTLE	100.0	Always-effective
subtle_author_pref	SUBTLE	99.0	Always-effective
context_confusion	DIRECT	95.9	Always-effective
subtle_compliance	SUBTLE	19.4	Length-sensitive
subtle_citation	SUBTLE	18.4	Length-sensitive
instruction_leak	DIRECT	3.1	Length-sensitive
flattery_social	DIRECT	2.0	Length-sensitive
direct_override	DIRECT	21.4	Uniformly weak
role_manipulation	DIRECT	14.3	Uniformly weak

Table 4: ASR (%) by injection position and attack category on GPT-4.1-MINI. The end-of-document position (100%) dramatically boosts direct attack success.

CATEGORY	INJECTION POSITION						
	0%	10%	25%	50%	75%	90%	100%
DIRECT	20.0	21.4	21.4	25.7	22.9	18.6	61.4
SUBTLE	70.0	70.0	65.7	64.3	68.6	65.7	67.1
Overall	45.0	45.7	43.6	45.0	45.7	42.1	64.3