

category.




Figure 17. (left) Comparing SOTA attacks from HarmBench and MSJ on Claude 2.0: We find that MSJ is more effective than the PAIR (Chao et al., 2023) and TAP (Mehrotra et al., 2023), ART (Perez et al., 2022a). PAP (Schulhoff et al., 2023) and the human attacker baselines are not shown, since they don’t jailbreak Claude 2.0 on any of the 200 HarmBench categories. **(right) Comparing MSJ variants on Claude 2.0:** We observe that both composing MSJ with other jailbreaks (Compositional MSJ), and sampling the in-context demonstrations from the same HarmBench category (referred to as “Same Category” in the legend) boost attack success rate. Note that MSJ still gets stronger as the context length grows and might not need any further bells and whistles given larger context windows.